

TCP vs UDP



Ashish Pratap Singh



⌚ 5 min read

When you send a WhatsApp message, stream a YouTube video, or play an online game, your device relies on a set of rules to move data across the internet.

At the heart of this process are two fundamental protocols: **TCP** and **UDP**. They are the invisible couriers that decide *how* your data is packaged and sent.

Both TCP (Transmission Control Protocol) and UDP (User Datagram Protocol) operate at the **Transport Layer (Layer 4)** of the OSI model. Their job is to take data from an application on one device and deliver it to an application on another.

Choosing between them is a critical system design decision.

- Do you need a courier that guarantees every single piece

of your message arrives in perfect order, even if it takes a little longer?

- Or do you need one that delivers as fast as possible, even if a few pieces get lost along the way?

This choice is a constant balancing act between **speed**, **reliability**, and **resource usage**.

1. The Role of the Transport Layer

Before diving into TCP and UDP, let's clarify what the **Transport Layer** does. While the **Network Layer (IP)** gets data packets from one computer to another (host-to-host), the Transport Layer gets data from **one application to another (process-to-process)**.

Its key responsibilities include:

- **Segmentation and Reassembly:** Breaking large chunks of application data into smaller packets on the sender's side and reassembling them on the receiver's side.
- **Delivery Model:** Providing either a **reliable delivery** service or a **best-effort** transmission.
- **Flow and Congestion Control:** Managing the rate of data

transmission to avoid overwhelming the receiver or the network itself.

This is where TCP and UDP diverge. **TCP focuses on providing a highly reliable, ordered stream of data, while UDP prioritizes speed and simplicity with a best-effort approach.**

2. What Is TCP (Transmission Control Protocol)?

TCP is a **connection-oriented** protocol that guarantees the reliable delivery of data in the correct order. It's the meticulous, trustworthy courier of the internet. Before sending any data, TCP establishes a formal connection between the client and server.

Key Characteristics:

- **Connection Establishment:** TCP uses a **three-way handshake** (SYN, SYN-ACK, ACK) to establish a reliable connection. This initial negotiation adds a bit of latency but ensures both sides are ready to communicate.
- **Reliability:** Every packet sent is tracked with a sequence number. The receiver sends acknowledgments (ACKs)

for packets it receives. If the sender doesn't get an ACK within a certain time, it **retransmits** the lost packet.

- **Ordered Delivery:** Sequence numbers also ensure that packets are reassembled in the correct order at the destination, even if they arrive out of order.
- **Flow & Congestion Control:** TCP uses a "sliding window" mechanism to prevent the sender from overwhelming the receiver (**flow control**). It also intelligently slows down transmission when it detects network congestion (**congestion control**).

How TCP Works

1. Connection Establishment (Three-Way Handshake)

Before any data exchange can happen, TCP establishes a connection between the client and server using a three-step handshake. This ensures that both sides are ready and can communicate reliably.



<--- SYN-ACK (I'm ready, are you?)

----- ACK (Yes, I am. Let's talk)

- **SYN (Synchronize):** The client initiates the connection by sending a SYN message, signaling its intent to start communication and sharing its initial sequence number (ISN).
- **SYN-ACK (Synchronize-Acknowledge):** The server receives the SYN, reserves resources for the connection, and responds with a SYN-ACK to acknowledge the request and share its own ISN.
- **ACK (Acknowledge):** The client sends an ACK to confirm receipt of the SYN-ACK, completing the handshake.

At this point, both sides have synchronized sequence numbers, and a reliable communication channel is established.

2. Data Transfer

Once the connection is established, TCP begins transmitting data. The protocol guarantees that all packets (segments) are delivered **in order, without duplication, and without loss**.


- **Segmentation:** Large data is broken into smaller, manageable chunks called *segments*.


- **Acknowledgments (ACKs):** Each segment sent is acknowledged by the receiver. If the sender doesn't receive an ACK within a set timeout, it retransmits the segment.
- **Sliding Window:** TCP uses a window-based flow control mechanism that allows multiple packets to be "in flight" before requiring an acknowledgment, improving throughput.
- **Error Detection:** Each segment includes a checksum to detect corruption in transit. Corrupted packets are discarded and resent.

In essence, TCP behaves like a reliable courier service: it delivers every packet, ensures they arrive in the right order, and resends any that get lost.

3. Connection Termination

When communication is complete, TCP closes the connection gracefully through a four-step termination process. This ensures both client and server agree that the session is finished.





```
sequenceDiagram
    participant Client
    participant Server
    Client->>Server: FIN (I'll close too)
    activate Client
    Client-->>Server: ACK (Got it)
    deactivate Client
    Client->>Server: FIN (I'll close too)
    activate Client
    Client-->>Server: ACK (Got it)
    deactivate Client
```

- **FIN (Finish):** The client initiates termination, signaling it has no more data to send.
- **ACK:** The server acknowledges the FIN, allowing any remaining data to be processed.
- **FIN:** The server sends its own FIN to close its side of the connection.
- **ACK:** The client confirms, and both sides release their resources.

This **graceful shutdown** ensures that no data is lost and both parties are aware the connection is closed.

3. What Is UDP (User Datagram Protocol)?

UDP is a **connectionless** protocol that sends data without establishing a formal connection. It operates on a "fire-and-forget" principle. It's the fast, no-frills courier of the inter-

net.

Key Characteristics:

- **No Handshake:** Packets (called datagrams) are sent immediately without any prior negotiation. This significantly reduces initial latency.
- **No Acknowledgments:** UDP doesn't wait for ACKs and doesn't retransmit lost packets. If a packet is dropped, it's gone for good.
- **No Ordering:** There's no guarantee that packets will arrive in the order they were sent.
- **Lightweight:** The UDP header is much smaller (8 bytes) than the TCP header (20+ bytes), meaning less overhead per packet.

How UDP Works


UDP doesn't perform a handshake like TCP. There's no set-up phase, no exchange of sequence numbers, and no acknowledgment of connection readiness.

When an application wants to send data:

- It creates a **datagram**, attaches a destination IP and port, and sends it directly to the network.
- The receiving application (if it's listening on that port)

simply processes the incoming data.

This simplicity eliminates the connection overhead, allowing UDP to deliver data almost instantaneously.



UDP treats every packet (called a **datagram**) as an independent message. Each datagram carries its own header information and is routed individually through the network.

This makes UDP ideal for real-time communication where dropping a few packets is acceptable but delays are not.

4. Head-to-Head Comparison: TCP vs UDP

This table summarizes the core differences:

Feature	TCP (Transmission Control Protocol)	UDP (User Datagram Protocol)
Type	Connection-oriented	Connectionless
Reliability	Reliable (acknowledged delivery)	Unreliable (no ACKs, "best-effort")
Order	Guaranteed (sequenced packets)	Not guaranteed
Speed	Slower due to hand-shakes and overhead	Faster, minimal overhead
Packet Loss Handling	Retransmits lost packets	Ignores lost packets
Flow & Congestion Control	Yes	No
Use Cases	Web, Email, File Transfer, APIs, Databases	Streaming, Gaming, VoIP, DNS

In essence, **TCP prioritizes correctness, while UDP prioritizes timeliness.**

5. Modern Innovations

The classic TCP vs. UDP debate has evolved. Modern protocols aim to get the best of both worlds.

QUIC (Quick UDP Internet Connections)

Developed by Google and now the basis for **HTTP/3**, QUIC is a game-changer.

- It runs over **UDP** to avoid the initial latency of the TCP handshake.
- It builds reliability, congestion control, and stream management directly into its own layer.
- It features built-in, mandatory encryption (TLS 1.3).
- It supports **multiplexing**, where multiple data streams can be sent over a single connection without one blocking the others.

Essentially, **QUIC provides many of TCP's benefits with the speed of UDP.**

Some applications also build their own **custom reliability layers** over UDP. For example, a multiplayer game might use UDP for player movement but implement a simple ACK system for critical events like "player used a special ability."

Master System Design

Progress

11/130 chapters

Q Search topics...



Introduction

2/3 ▾

Core Concepts

6/8 ▾

Networking

3/9 ▲



Get Premium

Subscribe to unlock full access to all premium content

Subscribe Now

Reading Progress

100%

On this page

1. The Role of the Transport Layer

2. What Is TCP (Transmission Control Protocol)?...

- OSI Model
- IP Address
- Domain Name System (DNS)
- Proxy vs Reverse Proxy
- HTTP/HTTPS

TCP vs UDP

- Load Balancers
- Load Balancing Algorithms
- Checksums

API Fundamentals

0/10 ▾

Databases & Storage

0/12 ▾

Database Scaling Techniques

0/8 ▾

Caching

0/6 ▾

Asynchronous Communications

0/4 ▾

Tradeoffs

0/9 ▾

Throughput

- TCP:** Its congestion control mechanism can throttle throughput to adapt to network conditions. This is good for the overall health of the internet but can limit the maximum speed of a single connection.
- UDP:** Can achieve higher throughput as it doesn't slow down for congestion. The risk is that sending too fast can lead to more packet loss.

Reliability vs. Efficiency

- TCP:** Provides reliability "for free" at the transport layer

6. Choosing the Right Protocol

When deciding between TCP and UDP, consider these factors:

Latency

- TCP:** Introduces latency upfront with its three-way handshake. Every acknowledgment and potential retransmission adds to the round-trip time (RTT).
- UDP:** Has near-zero initial latency. Data is sent immediately. This is a huge advantage for time-sensitive applications.

Throughput

How TCP Works

3. What Is UDP (User Datagram Protocol)?

How UDP Works

4. Head-to-Head Comparison: TCP vs UDP

5. Modern Innovations

6. Choosing the Right Protocol

7. Real-World Use Cases



but uses more CPU, memory, and bandwidth to manage connections and state.

- **UDP**: Is highly efficient but pushes the responsibility of handling packet loss and ordering up to the application layer.

Security

- **TCP**: Is the foundation for **TLS (Transport Layer Security)**, which powers HTTPS. The reliable, ordered nature of TCP is a prerequisite for how TLS works.
- **UDP**: Does not have a built-in mechanism for security. Applications must use a protocol like **DTLS (Datagram TLS)** to add a layer of encryption.

Here's a simple decision framework for your next system design interview or project:

If your application requires...	Then use...
Guaranteed & Ordered Delivery	TCP
The Lowest Possible Latency	UDP (or QUIC)
Tolerance for some Packet Loss	UDP
High Accuracy and Data Integrity	TCP
Real-Time Communication	UDP
Built-in Flow/Congestion Control	TCP

Many large-scale systems use a **hybrid approach**:

- **TCP** for critical user authentication and API calls.
- **UDP** for sending non-essential telemetry and metrics.
- **QUIC/HTTP/3** for user-facing web content and real-time streaming.

7. Real-World Use Cases

The choice between TCP and UDP is dictated entirely by the application's requirements.

TCP-Based Systems:

- **Web Traffic (HTTP/HTTPS):** When you load a web page, every single byte of HTML, CSS, and JavaScript must arrive correctly and in order. A missing packet could break the page layout or functionality.
- **Databases:** A database query or transaction must be 100% complete and correct. You can't afford to have a partially updated record due to a lost packet.
- **Email (SMTP, IMAP):** Just like a web page, an email message must be delivered in its entirety to be readable.

UDP-Based Systems:

- **Video Streaming (YouTube, Netflix):** If a single frame of a video is dropped, the user might see a tiny glitch or nothing at all. It's better to skip that frame and show the next one in real-time than to pause the entire stream to wait for a retransmission.
- **Online Gaming:** In a fast-paced game, receiving slightly outdated information is better than receiving perfect information too late. The game client can often interpolate or predict the missing data.
- **Voice over IP (VoIP - Zoom, WhatsApp Calls):** A small, momentary audio dropout is preferable to a long, jarring delay while the system tries to recover a lost packet.

