

# Distributed Transactions



Ashish Pratap Singh

11 1 1

4 min read



Get Premium

Subscribe to unlock full access to all premium content

Subscribe Now

A **distributed transaction** is a transaction that spans multiple, physically or logically separate databases or systems. Unlike a traditional transaction confined to a single system, a distributed transaction ensures that a set of operations across different systems either all succeed (commit) or all fail (rollback), maintaining data consistency.

Imagine you're at a busy restaurant that uses separate systems for orders, payments, and kitchen operations. A distributed transaction would ensure that when a customer places an order, the order is recorded, the payment is processed, and the kitchen is notified—all together.

If one part fails (say, the payment system), the entire transaction is canceled, preventing a scenario where the order is taken but payment isn't processed.

## 1. Why Do We Need Distributed Transactions?

In today's distributed architectures, particularly with the rise of microservices and multi-database systems, data is no longer stored in a single centralized system. Instead, it is spread across multiple services and locations.

Distributed transactions are critical because they:

- **Maintain Consistency:** They ensure that all parts of a transaction are completed successfully, keeping the system in a consistent state.
- **Prevent Partial Failures:** By committing or rolling back all related operations together, they avoid scenarios where some systems are updated while others are not.
- **Support Business Integrity:** For processes like financial transactions, order processing, or inventory management, distributed transactions guarantee that business rules are followed correctly across the board.

## 2. Key Concepts: ACID in a Distributed World

To understand distributed transactions, it's important to revisit the **ACID** properties that define a reliable transaction in a database:

Reading Progress

5%

On this page

1. Why Do We Need Distributed Transactions?
2. Key Concepts: ACID in a Distributed World
3. How Distributed Transactions Work
4. Challenges and Trade-Offs
5. Best Practices
6. Real-World Use Cases
7. Conclusion

- **Atomicity:** All parts of the transaction succeed or fail as one unit.
- **Consistency:** The system remains in a valid state before and after the transaction.
- **Isolation:** Transactions do not interfere with each other.
- **Durability:** Once committed, the transaction's changes persist, even in the event of a system failure.

Distributed transactions extend these ACID properties across multiple systems, which adds significant complexity.

## 3. How Distributed Transactions Work

### Two-Phase Commit (2PC)

The **Two-Phase Commit (2PC)** protocol is one of the most common methods to achieve distributed transactions.

#### Phase 1: Prepare (Voting Phase)

- **Coordinator Node:** A designated coordinator sends a prepare message to all participant nodes, asking if they can commit the transaction.
- **Participants:** Each participant executes the transaction locally and replies with a vote: "Yes" if the transaction can be committed, or "No" if it fails.

#### Phase 2: Commit or Abort

- **If All Votes are "Yes":** The coordinator sends a commit command to all participants, and each node commits the transaction.
- **If Any Vote is "No":** The coordinator sends an abort command, and all nodes roll back the transaction.

### Three-Phase Commit (3PC)

The **Three-Phase Commit (3PC)** protocol introduces an additional phase to reduce the risk of blocking, which can occur in 2PC if the coordinator fails. The phases are:

1. **CanCommit:** Similar to the prepare phase in 2PC.
2. **PreCommit:** The coordinator informs participants to prepare for commit, allowing a buffer phase.
3. **DoCommit:** Finally, the coordinator instructs participants to commit.

3PC provides a non-blocking mechanism under certain conditions, though it is more complex and less commonly used.

## 4. Challenges and Trade-Offs

## Challenges

- **Complexity:** Distributed transactions add significant complexity, especially with protocols like 2PC and 3PC.
- **Performance Overhead:** Coordination and multiple phases can introduce latency and reduce throughput.
- **Partial Failures:** Handling failures in a distributed environment is inherently challenging, requiring robust error detection and compensation mechanisms.
- **Scalability:** Traditional distributed transaction protocols might not scale well in highly distributed or high-volume systems.

## Trade-Offs

- **Consistency vs. Performance:** Ensuring strict consistency (using 2PC/3PC) can slow down the system, whereas approaches like sagas offer eventual consistency with better performance.
- **Complexity vs. Simplicity:** More complex protocols offer stronger guarantees but are harder to implement and maintain.

## 5. Best Practices

- **Minimize the Scope:** Limit the distributed transaction to only the necessary services. Use local transactions as much as possible.
- **Favor Sagas for Microservices:** In many modern architectures, the saga pattern is preferred over traditional distributed transactions for better scalability and fault tolerance.
- **Implement Robust Monitoring:** Use logging, tracing, and monitoring tools to track the progress of distributed transactions and quickly detect failures.
- **Design for Failure:** Assume that failures will occur and plan compensating actions accordingly.
- **Automate Testing:** Use simulations and chaos engineering to test how your system behaves under failure scenarios and during distributed transactions.

## 6. Real-World Use Cases

- **Financial Services:** Ensuring that transactions like fund transfers either complete fully or are fully rolled back, maintaining account consistency.
- **E-Commerce:** Coordinating order processing, payment, and inventory updates to ensure that an order is either fully processed or completely canceled.
- **Booking Systems:** Managing reservations across different services (e.g., flights, hotels, car rentals) where each step must succeed, or compensating actions must be taken.

- **Distributed Databases:** Maintaining consistency across multiple nodes or shards in a distributed database system.

## 7. Conclusion

Distributed transactions are critical for ensuring data consistency and integrity across multiple systems in today's complex, distributed environments. While traditional protocols like Two-Phase Commit and Three-Phase Commit provide strong guarantees, they come with significant complexity and performance overhead. Alternative approaches, such as the Saga Pattern, offer a more scalable solution with eventual consistency, which is often acceptable in microservices architectures.

By carefully weighing the trade-offs, implementing robust monitoring, and designing for failure, you can build distributed systems that maintain consistency without sacrificing performance. Whether you're working on financial systems, e-commerce platforms, or any distributed architecture, understanding and effectively implementing distributed transactions is key to building resilient and reliable systems.

[← Prev: Leader Election](#)[Take Notes](#)[Star](#)[Mark as Complete](#)[Ask AI](#)[Next: Gossip Protocol >](#)